

# Learning Adversarial Search algorithms

The Art of Unlosable Tic-Tac-Toe

---

Eric Han

June 7, 2024

# Introduction

---

# Experience

My *Singaporean* educational journey to CS, R&D:

- [2008] **Secondary School** — Informal learning; scripting, games
  - Interest in Computing: why & how computers work
- [2010] **Pioneer JC** — H2 Computing
  - Interest in research: A\*STAR IHPC Quest 2009 (Bronze) - K-Means
- [2018] **B.Com. NUS** — Com. Sci w Honors
  - A\*STAR Scholarship - Internships working on R&D projects
- [2024] **PhD. NUS** — AI/ML tackling scaling and robustness
  - First-author publications in AAAI, ICML.

Working experiences:

- [2018-2024] **Teaching Assistant/Graduate Tutor, NUS** — Teach UG
- [B.Com.] **Research Intern, A\*STAR IHPC** — ML Platform, Rec. Sys.

Support, teach ( $> 500$  contact hours), grade, manage/mentor tutors for:

- AI/Machine Learning
  - **CS2109s** Introduction to AI and Machine Learning
  - **CS3243** Introduction to Artificial Intelligence
- Software Engineering
  - **CS3217** Software Engineering on Modern Application Platforms
  - **CS3203** Software Engineering Project
  - **CS2030/CS2030S** Programming Methodology II

Skilled with Linux (also administration), Windows, macOS:

- **Programming Languages** — Python, C++, Java, Mojo...
- **Databases** — Firebase, SQL...
- **Typesetting / Presentation Tools** — LaTeX, Markdown (This slides!)
- **Tools/Platforms** — Git, Mlflow, Plotly, Slurm, GCP...

# Teaching Philosophy

Effective learning is driven by an *innate desire* to learn the subject rather than *need*:

1. **Creating a relaxed and safe environment** — Informal, casual, personal.
2. **Engaging students to facilitate learning in and after class** — Telegram, buddy
3. **Creating equal opportunities for all students to learn** — Reaching out

## Teaching Excellence (Tutorials/Recitation)

	2109	2109	3243	3243	3217	3203	3203	3203	3203	3203
<b>Score</b>	4.8	4.6	4.8	4.5	3.8	4.6	4.4	4.8	4.1	3.3
<b>Resp.</b>	36	13	25	39	6	13	16	18	20	3
<b>Nom.</b>	47%	30%	32%	31%	0%	61%	31%	61%	10%	33%

For the teaching position, I am interested to

- Focus on improving teaching quality
- Curriculum development/improvement
- Casual research
  - Mentor for Undergraduate Research / FYP
- Involved in consultancy/policy

## Mini-Lecture

---

## Recap on environment properties

- **Fully / Partially Observable:** Can the agent see?
- **Single / Multi-Agent:** How many agents?
- **Deterministic / Stochastic:** Is there randomness in transition?
- **Episodic / Sequential:** Is there dependence on previous action?
- **Static / Dynamic:** Can the environment change while the agent is thinking?
- **Discrete / Continuous:** Discretized or varying continuously?



# Recap on formulation

## Un/Informed Search (Path):

- State space
- Initial state
- Final state
- Action
- Transition

## Local Search (Goal):

- Initial state
- Transition
- Heuristic/Stopping criteria

**Motivation:** How can we win?

Ingredients needed to formulate a problem:


- **Initial state:** Starting configuration (representation)
- **Players:** Decision-makers within the game (2 players)
- **Actions:** Potential moves that the player can make
- **Transition:** Result of a move from a state
- **Terminal/Leaf test:** Checks if the game is over
- **Utility:** Reward for a terminal state and player

## Tic-tac-toe

2P childhood game where  $(P_O, P_X)$  players take turns drawing their symbols on a 3x3 grid. The winner is the first player to get 3 of his/her symbol in a row, col. or diag.


# Tic-tac-toe

2P childhood game where  $(P_O, P_X)$  players take turns drawing their symbols on a 3x3 grid. The winner is the first player to get 3 of his/her symbol in a row, col. or diag.



# Tic-tac-toe

2P childhood game where  $(P_O, P_X)$  players take turns drawing their symbols on a 3x3 grid. The winner is the first player to get 3 of his/her symbol in a row, col. or diag.



## Recap — Environment Properties

Fully Observable, 2 Agent, Deterministic, Squential, Static, Discrete

## Modeling Tic-tac-toe [Discussion]



2P childhood game where  $(P_O, P_X)$  players take turns drawing their symbols on a 3x3 grid. The winner is the first player to get 3 of his/her symbol in a row, col. or diag.

- **Initial state:**
- **Players:**
- **Actions:**
- **Transition:**
- **Terminal/Leaf test:**
- **Utility:**

# Modeling Tic-tac-toe

$$S_0 = \begin{array}{|c|c|c|} \hline 0 & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array} \xrightarrow{a_0=(1,X)} \begin{array}{|c|c|c|} \hline & X & \\ \hline & & \\ \hline & & \\ \hline \end{array} \rightarrow \dots \rightarrow \begin{array}{|c|c|c|} \hline O & X & X \\ \hline X & X & O \\ \hline O & O & \\ \hline \end{array} \xrightarrow{a_8=(8,X)} \begin{array}{|c|c|c|} \hline O & X & X \\ \hline X & X & O \\ \hline O & O & X \\ \hline \end{array} = S_9$$

- **Initial state:**  $S_0$ , 1D array of 9 elements:  $O, X, \emptyset$
- **Players:**  $P_X$ -max,  $P_O$ -min
- **$i$ -th Actions:**  $a_i = (c_i, y) : c_i \in [0, 8]$  where  $S_i[c_i] = \emptyset$ , symbol  $y \in X, O$
- **Transition:**  $T(S_i, a_i) = S_{i+1}$ , where  $S_{i+1}[j] = \begin{cases} y & \text{if } j = c_i \\ S_i[j] & \text{otherwise} \end{cases}$
- **Terminal/Leaf test:** Row, col. or diag having same symbols or no moves
- **Utility:**  $U(S_i, p)$  is 0 if draw, 1 if  $p$  wins,  $-1$  if  $p$  loses

# Modeling Tic-tac-toe

- **Initial state:**  $S_0$ , 1D array of 9 elements:  $O, X, \emptyset$
- **Players:**  $P_X$ -max,  $P_O$ -min
- **$i$ -th Actions:**  $a_i = (c_i, y) : c_i \in [0, 8]$  where  $S_i[c_i] = \emptyset$ , symbol  $y \in X, O$
- **Transition:**  $T(S_i, a_i) = S_{i+1}$ , where  $S_{i+1}[j] = \begin{cases} y & \text{if } j = c_i \\ S_i[j] & \text{otherwise} \end{cases}$
- **Terminal/Leaf test:** Row, col. or diag having same symbols or no moves
- **Utility:**  $U(S_i, p)$  is 0 if draw, 1 if  $p$  wins,  $-1$  if  $p$  loses

FAQ: Can I describe and not write math?

Yes, but it must be **clear**; ie. Able to translate into code without additional assumptions; you should (at min) describe how the state is represented.



# Modeling Tic-tac-toe in Python

- **Initial state:**  $S_0$ , 1D array of 9 elements:  $O, X, \emptyset$
- **$i$ -th Actions:**  $a_i = (c_i, y) : c_i \in [0, 8]$  where  $S_i[c_i] = \emptyset$ , symbol  $y \in X, O$

```
X,O,E = 'X','O','.'
```

```
SYMBOLS = {X,O}
```

```
WINNING_POS = [[0,1,2], [3,4,5], [6,7,8], [0,3,6], [1,4,7], [2,5,8], [0,4,8], [2,4,6]]
```

```
class TicTacToe(object):
```

```
    def __init__(self, Si=[ E ] * 9):
```

```
        self.Si = Si
```

```
        self.winner = E
```

```
    def actions(self):
```

```
        e_cis = [ ci for ci, Si_ci in enumerate(self.Si) if Si_ci == E ]
```

```
        return [ (ci, y) for y in SYMBOLS for ci in e_cis ]
```

## Zero-sum game

Zero-sum game is a game where one player gain is equals to another's loss, where the total utility of the game is the same/constant (ie. no improvement).

### Tic-tac-toe is zero-sum

- If  $P_X$  wins  $P_O$  loses:  $\sum U = 1 - 1 = 0$
- If  $P_O$  wins  $P_X$  loses:  $\sum U = 1 - 1 = 0$
- If  $P_O, P_X$  draws:  $\sum U = 0 + 0 = 0$


So, for Tic-tac-toe:  $U(S_i, X) = -U(S_i, O)$

**Intuition:** For nerds like me, if you played enough, you notice you keep getting draws.

#### Question

Can we come up with an algorithm to play Tic-tac-toe?


# Tic-tac-toe gametree



**Figure 1:** Gametree (R&N 3rd Ed) — Initial, Players, Actions, Transition, Terminal, Utility

# Minimax algorithm

**Intuition:** Simulate the game until the end with an imaginary optimal opponent.




I am player  $P_X$ , trying to find the best move at

O	X	
X	X	O
	O	

# Minimax algorithm


**Intuition:** Simulate the game until the end with an imaginary opponent.



We can fill in the utility values for the leaf nodes!

# Minimax algorithm


**Utility for X:**  $U(S_i, X)$  is 0 if draw, 1 if X wins,  $-1$  if X loses



We know we want the best action later, so we choose the best action (max) there!

# Minimax algorithm

For the best action chosen, we inherit its corresponding value!




We don't know how the  $P_O$  will play this move, so

- Assume that  $P_O$  wants to win and plays optimally like me.
- We imagine that  $P_O$  chooses the best action (min) there!

# Minimax algorithm

**Intuition:** Simulate the game until the end with an imaginary *optimal* opponent.



Now I can just pick the move that is the best (max value):

- All 3 moves would, at worse-case, end up in draws.




# Minimax algorithm

**Intuition:** Simulate the game until the end with an imaginary *optimal* opponent.

**function** MINIMAX-DECISION(*state*) **returns** *an action*

**return**  $\arg \max a \in \text{ACTIONS}(s) \text{ MIN-VALUE}(\text{RESULT}(state, a))$

# Minimax Tic-tac-toe example



Issues - Huge/Infinite Game Trees

# Alpha-Beta Pruning algorithm

Intuition

## Alpha-Beta Tic-tac-toe example









## Minimax 2048 example

# Minimax limitations

# Expectimax algorithm

## Expectimax example